

Assignment A3: Team 33

Lucas van Osenbruggen

Martin Miksik

Luca Mainardi

1 Agent Description

We implement an agent for the game *Competitive Sudoku*, in which two players must complete a Sudoku puzzle one turn at a time and receive points for completing regions (Kraak [2023]). Two important rules are: A) The agent cannot suggest *taboo moves*, which are listed by the game engine. Making such move, which would render the puzzle unsolvable, results in losing the game. B) If a player makes a move that makes the Sudoku unsolvable and the move is not yet marked as a taboo move, that move is added to the taboo list, and the turn passes to the opponent. We design two agents, based respectively on the *Minimax* algorithm and on a *Monte Carlo tree search (MCTS)* strategy.

1.1 Iterative Deepening Minimax Agent

The Minimax agent, which can be found in Listing 2, uses a search strategy based on the *Minimax* algorithm (Polak [1989]). This is implemented in the function `minimax`. This strategy tries to maximise the score of the player, taking the potential moves of the opponent into account by searching a *game tree* of all subsequent moves. Since the turn time is indeterminate, we cannot reliably search to arbitrary depth. For a board of size $N \times N$, there are $\mathcal{O}(N^3)$ initial moves possible for the first turn, and $\mathcal{O}(N^3!)$ possible games. To deal with this time constraint, our agent will always start by proposing a random legal move. After this, the agent searches the game tree until some iteratively increasing depth – a strategy which is called *iterative deepening* (Korf [1985]). This is implemented in the function `compute_best_move`. Searching until a limited depth requires an evaluation function of an incomplete game. To this end, we define the following evaluation score:

$$\Delta(s_{max}, s_{min}) = s_{max} - s_{min} \quad (1)$$

Where s_{max} and s_{min} represent the game scores at the maximum search depth of the maximising and minimising player respectively. This evaluation score is optimal if the agent can search the entire game tree. This iterative deepening strategy ensures that the agent will propose the best move it can within the time limit. Furthermore, given more time, the agent will find better moves.

To address the limited time and improve performance, we implement several heuristics:

Alpha beta pruning is a strategy in which branches of the search tree that cannot lead to a better score than found so far are not searched (Knuth and Moore [1975]). This heuristic reduces the search space, while still producing optimal moves.

Move ordering heuristics is used once per turn to sort the list of initial moves based on the number of regions they allow to complete. Moves that complete three regions have the highest priority and are placed at the front of the list, followed by those that complete two, those that complete one, and finally, those that complete no regions. Once the sorting of the moves has been completed, the first move in the list is proposed (the one that *likely* leads to a greater increase in score) and the Minimax search with iterative deepening




Figure 1: Examples of unsolvable Sudoku when the examined block row (blue region) has 3 filled cells (left), two filled cells (centre) and one filled cell (right). In all three cases the value 4 makes Sudoku impossible.

begins with a depth limit 2; that allows the agent to consider the opponent’s next move. The search starts from the moves that are presumably best, since the list of moves is sorted. This heuristic is implemented in the function `order_moves`.

Propose unsolvable moves heuristic consists of calculating a list of moves that make the Sudoku unsolvable (which we call *unsolvable moves*) and proposing one of them every time, the best score returned by the Minimax search is negative. The agent, thus, exploits the game rule that allows the player to skip a turn, deliberately proposing a move that makes the game unsolvable if all evaluated moves benefit the opponent. Furthermore, unsolvable moves are not evaluated by the *Minimax* to reduce the search space and save computation time.

The calculation of the list of unsolvable moves is done only in the *mid game*, i.e., when more than 20% and less than 85% of the cells on the board are full¹.

The calculation of the list is computationally expensive, therefore we save the result to disk. At each subsequent turn, the list is loaded from the disk and updated (e.g. removing illegal moves). If the list is empty, a new list of unsolvable moves is re-generated². Nonetheless, the computation time associated with unsolvable moves is still an issue, given that the duration of a turn is unknown. If there isn’t sufficient time to complete the list of unsolvable moves. Under these circumstances, the agent repeatedly attempts to calculate this list in subsequent turns. As a result, it never moves to the Minimax search. To address this, the agent includes a check at the beginning of each turn. It determines whether it previously started but failed to complete the list. If this is the case, the agent repetitively increases the required cell fill rate by 10%.

This heuristic is implemented in the function `compute_best_move`. The algorithm that determines whether a game state (with an $N \times N$ board) is impossible is outlined in Algorithm 1 and is implemented in the function `is_unsolvable`.

In short, we examine all the rows of the blocks (3 rows made up of three cells in the case of a 9x9 Sudoku with 3x3 blocks). Any of the values not present in the row (blue region in Figure 1) could invalidate the Sudoku. If one of these values is not present in the same block as the row considered (red region in Figure 1), but is present in each of the rows not considered and of the columns corresponding to the empty cells (green regions in Figure 1), then none exists legal cell for that value in the block, and therefore Sudoku is impossible.

It is important to note that not all impossible states are detected by this algorithm. For example, the same check that was done on the block rows could also be done on the columns. However, for our purposes full list of unsolvable moves is not needed and

¹Our testing showed that outside this range there are very rarely moves that make Sudoku unsolvable.

²As noted later our unsolvable moves list is not exhaustive

Algorithm 1 Impossible Sudoku. Function name correspond to the implementation in Listing 4

```

function is_impossible(board)
  for all block in the board do
    for all block_row in the block do
      for all value in range  $[1, N]$  not present in the block_row do
        if if value is not in the red region and value is in all green regions then
          return true
        end if
      end for
    end for
  end for
  return false

```

would further increase the computation time required. Additionally, our tests show that the described algorithm can find *at least one* unsolvable move at almost any point in the game.

1.2 Tree-saving MCTS Agent

The second agent uses a *Monte Carlo tree search (MCTS)* strategy (Swiechowski et al. [2021]). Contrary to the Minimax agent from Section 1.1, which is an exhaustive search method on a game tree, the MCTS agent plays complete games with random moves, i.e., *random playouts*, and thus randomly samples different paths in the game tree. Since MCTS completes entire games, games are evaluated by their final score Δ as follows:

$$\Delta(s_{max}, s_{min}) = \begin{cases} 1 & \text{if } s_{max} > s_{min} \\ 0 & \text{else} \end{cases} \quad (2)$$

Where s_{max} is the score of the maximising player and s_{min} of the minimising player. MCTS estimates the quality of different moves based on the average score of a given move and improves this estimate after every playout. Given a longer turn time, this agent is thus expected to find better moves.

Furthermore, the MCTS agent makes better use of the available turn time. This is best illustrated with a hypothetical example of a game with a 1 second turn time. We will assume that the board is such that the Minimax agent can search until a depth of 5 in this case. However, because it does not know this, the agent will first search until lower depths, and thus spend its time searching to depths 2 and 3 before losing the turn while searching depth 4. In this case, the Minimax agent searched too shallowly and also wasted time. MCTS on the other hand, will play 99 random playouts, proposing a better move each time, before losing the turn in the 100th playout. In this example MCTS did not do any duplicate work, nor is there a high cost of losing the turn at the wrong moment.

The implementation of the MCTS agent can be found in Listing 3. The agent follows the basic structure of Monte Carlo tree search, which is outlined in Algorithm 2 with the function names corresponding to those in the listing. Because of the difference in strategy, the implementation is quite different from the Minimax agent, although the generation of legal moves and updating of the state are similar. The agents do share some utility functions which are listed in Listing 4. These functions perform game-specific logic that is not related to the strategy of the agent. Unit tests have been written for these functions as well, to ensure their correctness.

In MCTS a *Monte Carlo game tree* describes the quality of different states by their cumulative score q and how often they were evaluated n . Our agent extends basic MCTS

by saving this data structure between turns. When the turn starts the agent loads the data structure and moves in the tree according to the move it proposed in the previous turn and the move performed by the other agent³. When these moves are not in the tree it creates a new root node. This allows the agent to reuse work the previous turn if the state evaluation tree is deeper than 2 levels.

Algorithm 2 Monte Carlo Tree Search of the best move in a given turn. Function names correspond to the implementation in Listing 3

```

function iterate(root)
  while turn lasts do
    leaf  $\leftarrow$  select(root)  $\triangleright$ Repeatedly select node with highest UCB value until a leaf
    leaf  $\leftarrow$  expansion(leaf)  $\triangleright$ Add new child nodes and select one under some conditions
    apply_move_on_node(leaf)  $\triangleright$ Update state by applying all moves in path
     $\Delta \leftarrow$  simulation(leaf)  $\triangleright$ Apply random moves until the game is over, then calculate the score (Equation 2)
    backpropagation(leaf,  $\Delta$ )  $\triangleright$ Update cumulative score and visited of nodes
    bestchild  $\leftarrow$  arg maxchild  $\in$  root.children ( $\frac{\text{child.q}}{\text{child.n}}$ )
    return bestchild  $\triangleright$ Propose move
  if number of turns is a multiple of 10 then
    save root recursively
  end if
end while

```

To summarise, we propose an MCTS agent that estimates the best moves by randomly sampling complete games from the state space and saving the estimates between turns.

2 Agents Analysis

In our evaluation of AI agents Minimax and MCTS, we focused on their performance relative to computational time constraints and varying game board configurations. We run the Minimax agent against the A2 agent, Greedy Player, and MCTS agent across different boards at 0.1, 0.5, 1, and 2-second intervals for 20 games each. This resulted in 12 distinct outcomes. We also run a similar set of experiments for the MCTS agent. Additionally, we test the performance of hyper-parameters for the MCTS agent. Finally, we do an analysis of the number of Minimax calls and MCTS iterations as a function of the time limit, to test the agents' ability to make the most of the time available and their robustness to different time conditions

2.1 The Minimax Agent Analysis

Performance Over Time Figure 2 illustrates how the agents' average win rates vary with the amount of computation time provided. We observed the following:

- Against the A2 agent, the Minimax agent outperforms A2 when the computation time is limited, suggesting that the Minimax heuristics improve performance. With more computation time the impact of heuristics disappears and the agents end up in a draw.

³An exception is a taboo move, which was proposed but not accepted and thus does not change the state.




Figure 2: Average performance of A3 agent over time

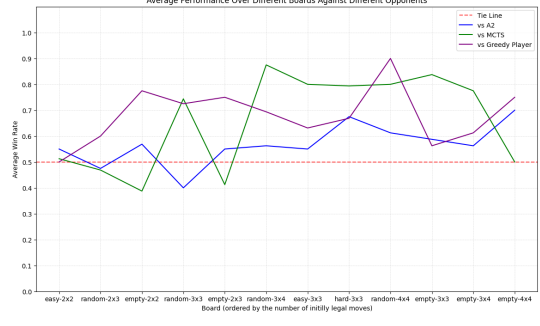


Figure 3: Average performance of A3 agent over different board configurations

- When facing the MCTS agent($C=2$)⁴ opponent, the Minimax agent starts with a high win rate which decreases as computation time increases. This suggests that MCTS may benefit more from additional computation time than A3. This is surprising since we would expect the MCTS to perform better with limited time. Minimax with enough time should provide *optimal* moves.
- Against the Greedy Player, the A3 agent's win rate increases over time, highlighting that the A3 agents can explore larger game trees and find *better* moves.

Performance Across Board Configurations Figure 3 presents the agents' win rates across different board setups' average overall times. We observed the following insights:

- Performance is generally inconsistent across different board configurations, suggesting that our A3 agent's implementation may not generalise well and thus its performance is board configuration dependent.
- The A3 agent seems to perform better on larger, empty boards (e.g., empty-3x4, empty-4x4) against the A2 agent. Indicating again that the A3 implementation better exploits the difference between *early*, and *late* game.
- The A3 agent generally outperforms all opponents (the win rate is above the tie line), with some exceptions, particularly on smaller boards. This indicates that the A3 agent performs better when the number of initial legal moves increases.

2.2 MCTS Agents Analysis

Performance Over Time Figure 4 showcases the fluctuation of the average win rates of the agents given different computation times. We observed the following:

- Against the Greedy Player, the MCTS agent shows an improvement in performance as the computation time extends, but in general, it does not win any matches.
- Competing with another MCTS, for time<0.5s the agents perform somewhat similarly. For time=1s, but with increased computation time decreasing the C hyper-parameter seems to have a positive effect on the performance.

Performance Across Board Configurations Figure 5 illustrates the agents' win rates over an average of all computation times across various board types. We observed the following:

⁴We chose to default to $C=2$ as per lecture slides

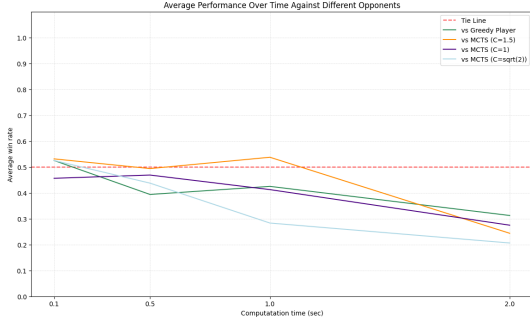


Figure 4: Average performance of MCTS (C=2) agent vs other agents (including varying C) over time

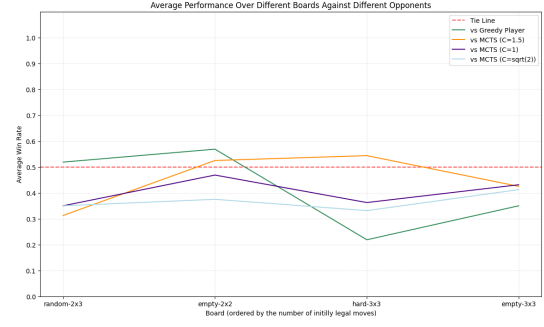


Figure 5: Average performance of MCTS (C=2) agent vs other agents (including varying C) over board configurations

2.3 Minimax Calls and Iterations of MCTS Analysis

Finally, an analysis was made of the difference in behaviour between Minimax and Monte Carlo Tree Search (MCTS), in particular examining the number of iterations as the time available to make a move varies. The number of calls of the Minimax algorithm and the number of MCTS iterations were measured by playing games on four boards (empty-2x2, random-2x3, hard-3x3, empty-3x3) with four different time limits (0.1s, 0.5s, 1s, 2s).

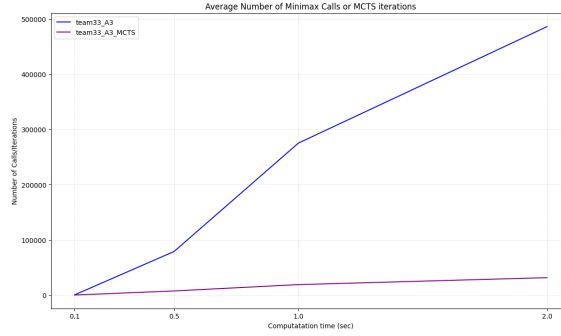


Figure 6: Average number of Minimax Calls and MCTS Iterations over time

The results in Figure 6 demonstrate that the number of calls of the Minimax function increases significantly as the time available for move selection increases, suggesting a tendency for the game tree to expand as the decision time extends. In contrast, the number of iterations of MCTS has a slight, almost constant slope, suggesting greater stability in carrying out iterations regardless of the time available. In other words, the analysis highlights a trade-off between deepening the decision tree and stability in the exploration of moves: Minimax adapts to the increase in available time by broadening its search, while MCTS maintains a more constant and risk-oriented strategy. Selective exploration, guided by Monte Carlo simulations.

3 Motivation & Reflection

In this section, we will reflect on several design decisions in the two agents and their strengths and weaknesses in the domain of Competitive Sudoku.

Incremental best moves improvements In Competitive Sudoku, the state space explodes with the board size. With a short turn time, agents often cannot explore all move

combinations and must search a subset of the state space, a challenge compounded by unknown time limits. Our agents tackle this by progressively improving the best move within their time limit. The Minimax agent searches the game tree to a certain depth, incrementally increasing this depth as time permits. This strategy leverages Competitive Sudoku’s scoring system, where completing regions earns intermediate scores, guiding the agent towards moves that offer short-term score gains. Conversely, the MCTS agent exploits the game’s characteristic of having relatively few turns in a game. It simulates numerous complete games to identify beneficial moves, continuously refining estimates after each simulated game.

We optimised the generation of new states by locally updating the boards based on the move, in the sense that applying a move at some position will only affect the row, column and block that position is in. This is helpful, as evaluating more states or games yields better results. For this reason, we measured this and the results are shown in Figure 6.

Data saving Both the MCTS and Minimax agents save data between turns to improve efficiency. The game state differs from the last turn only by the last two moves, making it unnecessary to recompute everything. The Minimax agent stores the computationally expensive list of *unsolvable moves*. The MCTS agent, maintaining a tree of potential best moves, updates its strategy by following the path of the last two moves in the saved game tree. This data retention in MCTS is particularly beneficial when the tree depth exceeds two layers. Longer turn times are hypothesised to enhance the MCTS agent’s performance. Additionally, the agent’s hyperparameter C , controlling the balance between exploration and exploitation, suggests that lower values (favouring exploitation) might lead to deeper trees and thus better performance. This was tested and the results are shown in Figure 4 and 5.

Intentional turn loss In Competitive Sudoku, some moves that would the Sudoku unsolvable are rejected, causing the agent to lose their turn. The Minimax agent exploits this by intentionally playing an unsolvable move when advantageous. For example, if the Minimax agent predicts that all explored moves will benefit the opponent, it may choose to forfeit its turn. This forces the opponent into a position where they might lose points. Additionally, the agent can avoid losing its turn when this would be disadvantageous. This strategic use of unsolvable moves adds a layer of complexity to the agent’s decision-making process.

Early vs late game The Minimax strategy is less effective in the early stage of the game when Equation 1 is often zero and thus not informative. This is partially mitigated by heuristically ordering moves. In contrast, it is quite effective in the late game when it can search most of the state space. MCTS makes better early game decisions by simulating complete games, but it is less effective at predicting short-term consequences, focusing less on immediate future nodes and under-penalising poor outcomes. Essentially, Minimax is nearsighted and MCTS farsighted. A prospective hybrid approach could combine MCTS’s early-game strength with Minimax’s late-game precision.

Symmetry Many solutions to a Sudoku are mirror images of others. As a result, many states result in the same score, and this property is stronger for states in nearly empty boards. Future work could consider an agent that is able to find this symmetry, allowing it to prune a large part of the search space.

References

- Donald E. Knuth and Ronald W. Moore. An analysis of alpha-beta pruning. *Artificial Intelligence*, 6(4):293–326, 1975. ISSN 0004-3702. doi:[https://doi.org/10.1016/0004-3702\(75\)90019-3](https://doi.org/10.1016/0004-3702(75)90019-3). URL <https://www.sciencedirect.com/science/article/pii/0004370275900193>.
- Richard E. Korf. Depth-first iterative-deepening: An optimal admissible tree search. *Artificial Intelligence*, 27(1):97–109, 1985. ISSN 0004-3702. doi:[https://doi.org/10.1016/0004-3702\(85\)90084-0](https://doi.org/10.1016/0004-3702(85)90084-0). URL <https://www.sciencedirect.com/science/article/pii/0004370285900840>.
- Thomas Krak. Competitive sudoku: Rules, 2023. URL https://canvas.tue.nl/courses/25392/files/5120538?module_item_id=522484. Accessed: 2023-11-29.
- E. Polak. *Basics of Minimax Algorithms*, pages 343–369. Springer US, Boston, MA, 1989. ISBN 978-1-4757-6019-4. doi:10.1007/978-1-4757-6019-4_20. URL https://doi.org/10.1007/978-1-4757-6019-4_20.
- Maciej Swiechowski, Konrad Godlewski, Bartosz Sawicki, and Jacek Mandziuk. Monte carlo tree search: A review of recent modifications and applications. *CoRR*, abs/2103.04931, 2021. URL <https://arxiv.org/abs/2103.04931>.

Python files

In this appendix, you must include all your Python files, so that the reviewers can evaluate your code. So this must at least be your `sudokuai.py` file, but might include any support files. Make sure that the file name (and path) is explicitly mentioned, so that it is clear how the code is used. Avoid too-long code lines that need to be broken over multiple lines here.

Code Listing 1: Heuristics Minimax agent with heuristics search from Assignment A2. `team33_A2/sudokuai.py`.

```
1  """Competitive Sudoku AI.  
  
3  Adapted from /naive_player/sudokuai.py  
  
5  A1: iterative deepening minimax search with alpha beta pruning.  
6  A2: heuristic search.  
7  """  
  
9  import os  
10 from random import shuffle  
  
12 # from numpy import full  
  
14 # Import types and libraries  
15 from competitive_sudoku.sudoku import GameState, Move, SudokuBoard  
16 from competitive_sudoku.sudokuai import SudokuAI  
  
18 from .utils import block_range # Game specific logic  
19 from .utils import block_index, calculate_move_score, is_illegal , next_player  
  
22 class SudokuAI(SudokuAI):  
23     """  
24     Sudoku AI agent that computes a move for a given sudoku configuration .  
25     """  
  
27     def __init__(self):  
28         super().__init__()  
29         self . transposition_table = []  
  
31     def update_legal( self , game_state: GameState, move: Move, moves: dict) -> list :  
32         """Update which moves are legal in the recursive minimax search.  
  
34         :param move: The move to be evaluated  
35         :param moves: A dictionary containing the initial set of moves, whether they  
                        are still legal and other properties of the  
                        moves.  
36         initial : list of initially legal moves, used as subset to avoid iterating  
                        over all moves  
37         legal : numpy array of shape (board_size, board_size, board_size + 1) where  
                        legal[i, j, k] is True if Move(i, j, k) is  
                        legal  
38         count: Counter for legal moves, avoid repeated iteration
```

```

39         :return: A list of moves that were invalidated by the move.
40         """
41         _moves_invalidated = []
42         for row in range(game_state.board.board_height()):
43             # move invalidates another move if not already illegal , avoid double
44                 counting
45             if moves["legal"][row, move.j, move.value]:
46                 _moves_invalidated.append(Move(row, move.j, move.value))
47                 moves["legal"][row, move.j, move.value] = False # Set legal status
48         for column in range(game_state.board.board_width()):
49             if moves["legal"][move.i, column, move.value]:
50                 _moves_invalidated.append(Move(move.i, column, move.value))
51                 moves["legal"][move.i, column, move.value] = False
52         for row, col in block_range(row=move.i, col=move.j, board=game_state.board):
53             if moves["legal"][row, col, move.value]:
54                 _moves_invalidated.append(Move(row, col, move.value))
55                 moves["legal"][row, col, move.value] = False
56         moves["count"] -= len(_moves_invalidated)
57         return _moves_invalidated
58
59     def minimax(
60         self,
61         game_state: GameState,
62         move: Move,
63         moves: dict,
64         current_player: int,
65         maximizing_player: int,
66         depth: int,
67         *,
68         alpha: float = float("-inf"),
69         beta: float = float("inf"),
70     ) -> float:
71         """Returns the score of a given move.
72
73         Minimax search that considers the perspectives of two players .
74         Searches until some depth before returning the score .
75         Uses alpha beta pruning to avoid searching branches which cannot lead to
76             better results .
77
78         :param game_state: The current game state. Describes the board, scores , taboo
79             moves and move history.
80         :param move: The move to be evaluated
81         :param moves: A dictionary containing the initial set of moves, whether they
82             are still legal and other properties of the
83             moves.
84             initial : list of initially legal moves, used as subset to avoid iterating
85                 over all moves
86             legal : numpy array of shape (board_size, board_size, board_size + 1) where
87                 legal[i, j, k] is True if Move(i, j, k) is
88                 legal
89             count: Counter for the number of legal moves.

```

```

82         free: shows per region (row, col or block) what the number of free squares
               is, used in heuristics .
83     :param current_player: The player who's turn it is. Will be the same
                           throughout the turn. (0 or 1 for first or
                           second player)
84     :param maximizing_player: The player who's score is to be maximised. (0 or 1
                           for first or second player)
85     :param depth: The maximum depth to search before returning the score.
86     :param alpha: The highest so far value for alpha beta pruning. ( initially  -inf
                   )
87     :param beta: The lowest so far value for alpha beta pruning. ( initially  inf)
88     :return: The score of the move (higher is better for maximizing player, lower
               is better for minimizing player)
89     """
90     block_indices = block_index(move.i, move.j, game_state.board)

91
92     # Apply move, update resulting scores
93     # Update legal moves and count newly invalidated moves
94     game_state.board.put(move.i, move.j, move.value)
95     _score_achieved = calculate_move_score(game_state, move)
96     game_state.scores[current_player] += _score_achieved
97     _moves_invalidated = self.update_legal(game_state, move, moves)
98     # switches perspective to other player
99     current_player = next_player(current_player)

100
101     # Update properties used in heuristics
102     moves["free"]["row"][move.i] -= 1
103     moves["free"]["col"][move.j] -= 1
104     moves["free"]["block"][block_indices[0]][block_indices[1]] -= 1

105
106     # Search until game is finished or maximum depth is reached
107     if depth == 0 or moves["count"] == 0:
108         # evaluate the current board
109         best_value = self.evaluate_state(
110             maximizing_player, current_player, game_state, moves["free"]
111         )
112     else:
113         if maximizing_player == current_player: # maximising player
114             best_value = float("-inf")
115             for try_move in moves["initial "]:
116                 if not moves["legal"][try_move.i, try_move.j, try_move.value]:
117                     continue
118                 # Recurse and find value up to some depth
119                 value = self.minimax(
120                     game_state,
121                     try_move,
122                     moves,
123                     current_player,
124                     maximizing_player,
125                     depth - 1,
126                     alpha=alpha,

```

```

127         beta=beta,
128     )
129     best_value = max(best_value, value)
130     alpha = max(alpha, best_value)
131     if beta < alpha:
132         break
133 else: # minimising player
134     best_value = float("inf")
135     for try_move in moves["initial "]:
136         if not moves["legal"][try_move.i, try_move.j, try_move.value]:
137             continue
138         # Recurse and find value up to some depth
139         value = self.minimax(
140             game_state,
141             try_move,
142             moves,
143             current_player,
144             maximizing_player,
145             depth - 1,
146             alpha=alpha,
147             beta=beta,
148         )
149         best_value = min(best_value, value)
150         beta = min(beta, best_value)
151         if beta < alpha:
152             break

153
154     # Undo move and its effects
155     current_player = next_player(current_player)
156     moves["free"]["row"][move.i] += 1
157     moves["free"]["col"][move.j] += 1
158     moves["free"]["block"][block_indices[0]][block_indices[1]] += 1
159     moves["count"] += len(_moves_invalidated)
160     for inv_move in _moves_invalidated:
161         moves["legal"][inv_move.i, inv_move.j, inv_move.value] = True
162     game_state.scores[current_player] -= _score_achieved
163     game_state.board.put(move.i, move.j, 0)

164
165     # Recursion result
166     return best_value

167
168 def evaluate_state(
169     self,
170     maximizing_player: int,
171     current_player: int,
172     game_state: GameState,
173     free: dict,
174 ) -> float:
175     """Heuristic evaluation of the current game state.

176
177     Is used by minimax to evaluate the current game state which is most often not

```

```

                                a complete game.
178     Base score is the difference between the scores of the two players since
                                maximizing this will result in a win.
179     Additional heuristic contributions are made for early game where score is
                                often 0.

181     :param maximizing_player: The player who's score is to be maximised. (0 or 1
                                for first or second player)
182     :param game_state: The current game state. Describes the board, scores, taboo
                                moves and move history.
183     :param free: The number of free squares per region.
184     :return: The score of the game state (higher is better for maximizing player,
                                and vice versa)
185     """
186     score = (
187         game_state.scores[maximizing_player]
188         - game_state.scores[next_player(maximizing_player)]
189     )

191     if not os.environ.get("not_prefer_more_empty"):
192         # Scale to avoid this additional heuristic dominating the score
193         # Will result in an early game strategy avoiding a filled field.
194         # Positive contribution, our player will thus prefer less filled fields.
195         if current_player == maximizing_player:
196             score += 0.1 * self.prefer_empty_regions(game_state, free)

198     return score

200 def prefer_empty_regions(self, game_state: GameState, free: dict):
201     """Heuristic for evaluating a state.
202
203     Prefer less filled out regions by counting the number of free moves
204     Normalising using the the maximum number of squares in a region
205
206     :param maximizing_player: The player who's score is to be maximised. (0 or 1
                                for first or second player)
207     :param game_state: The current game state. Describes the board, scores, taboo
                                moves and move history.
208     :param free: The number of free squares per region.
209     :return: The heuristic score, will be 1 for a completely empty field.
210     """
211     board_size = game_state.board.board_height()
212     num_blocks = board_size / game_state.board.region_height()
213     score = sum([count / board_size for count in free["row"]]) / board_size / 3
214     score += sum([count / board_size for count in free["col"]]) / board_size / 3
215     score += (
216         sum(
217             sum([count / board_size for count in block]) / num_blocks
218             for block in free["block"]
219         )
220         / num_blocks

```

```

221         / 3
222     )
223     return score

225     def find_initial_moves ( self , game_state: GameState) -> list[Move]:
226         """
227         Find all possible moves for a given state. This is copy of method used in A1
228                                     used for benchmarking purposes.
229
230         @param game_state: GameState
231         @return: list of moves
232         """
233         board_size = game_state.board.board_width()
234
235         # Generate possible moves
236         initial_moves = [
237             Move(i, j, value)
238             for i in range(board_size)
239             for j in range(board_size)
240             for value in range(1, board_size + 1)
241             if not is_illegal (move=Move(i, j, value), state=game_state)
242         ]
243
244         # Shuffle moves to be less predictable
245         shuffle ( initial_moves )
246
247         return initial_moves
248
249     def find_initial_moves_heuristics ( self , state : GameState) -> list[Move]:
250         """
251         Find all possible moves for a given state and order them by priority . The
252                                     priority is determined by the
253                                     number of possible moves for a cell . Cells with fewer possible moves are
254                                     prioritised because they are
255                                     more likely to result in a completed row, column, or block.
256                                     Any cell with 2 possible moves is *de*-prioritised , because that means the
257                                     opponent could complete a row, column
258                                     or block next turn.
259
260         @param state: GameState
261         @return: ordered list of moves
262         """
263         size = state.board.board_width()
264
265         # Store moves in a dictionary with the number of possible moves for that cell
266                                     as key
267         priority_dict = dict([(key, []) for key in range(0, size + 1)])
268
269         for i in range(size):
270             for j in range(size):
271                 possible_moves_for_cell = []
272                 for value in range(1, size + 1):
273                     move_candidate = Move(i, j, value)

```

```

267         if not is_illegal (move=move_candidate, state=state):
268             possible_moves_for_cell.append(move_candidate)

270         priority_dict [len(possible_moves_for_cell)].extend(
                possible_moves_for_cell)

272     key_order = sorted( priority_dict . keys())

274     # Prioritise cells that have 2 possible moves,
275     # because the opponent could complete a row, column or block
276     if len(key_order) > 2:
277         key_order.pop(2)
278         key_order.append(2)

280     # Return list of moves in order of search priority
281     return [move for key in key_order for move in priority_dict [key]]

283 def compute_best_move(self, game_state: GameState) -> None:
284     """Computes the best move for the agent and proposes it .

286     Will initially propose a random move, then evaluate different moves based on
287     minimax search.

289     Since the turn time is not known it will propose the best move found so far by
290     iteratively deepening the search .

291     :param game_state: The current game state. Describes the board, scores and
292     move history.
293     """

294     board_size = game_state.board.board_height()
295     num_blocks = board_size // game_state.board.region_height()

296     # print(f"Avoid 2 moves: {not os.environ.get('not_avoid_2_moves')}")
297     # print(f"Prefer empty regions: {not os.environ.get('not_prefer_more_empty')}")
298     # print(f"")

299     # Generate possible moves
300     initial_moves = (
301         self . find_initial_moves_heuristics (game_state)
302         if not os.environ.get("not_avoid_2_moves")
303         else self . find_initial_moves (game_state)
304     )

305     # Move cache
306     moves = {}

307     # List of initially legal moves, used as subset to avoid iterating over all
308     moves

309     moves["initial"] = initial_moves

310     # Propose a random move (initial, avoid timeout)
311     self . propose_move(moves["initial"][0])

```

```

313     # Avoid repeated regeneration of legal moves by tracking their status
314     from numpy import full
315     moves["legal"] = full(
316         shape=(board_size, board_size, board_size + 1),
317         dtype=bool,
318         fill_value=False,
319     )
320     moves["count"] = len(moves["initial"])
321     for move in moves["initial"]:
322         moves["legal"][move.i, move.j, move.value] = True

324     # Track some properties of the game to be used in statistical search
325     # Count how many free squares there are per region
326     # TODO do this above when generating moves?
327     moves["free"] = {
328         "row": [
329             len(set(((m.i, m.j) for m in moves["initial"] if m.i == y)))
330             for y in range(board_size)
331         ],
332         "col": [
333             len(set(((m.i, m.j) for m in moves["initial"] if m.j == x)))
334             for x in range(board_size)
335         ],
336         "block": [
337             [
338                 len(
339                     set(
340                         (
341                             (m.i, m.j)
342                             for m in moves["initial"]
343                             if block_index(m.i, m.j, game_state.board) == (i, j)
344                         )
345                     )
346                 )
347                 for j in range(num_blocks)
348             ]
349             for i in range(num_blocks)
350         ],
351     }

353     # Iteratively increase the search depth of minimax
354     best_move: tuple[Move, float] | None = None
355     for depth_limit in range(1, len(moves["initial"]), 1):
356         # evaluate different moves based on minimax
357         for move in moves["initial"]:
358             # player_index is 0 or 1 (first or second player)
359             player_index = game_state.current_player() - 1

361             value = self.minimax(
362                 game_state,

```



```

363         move,
364         moves,
365         player_index, # Current player is self
366         player_index, # Maximising own score
367         depth_limit,
368     )

370     if best_move is None or value > best_move[1]:
371         best_move = (move, value)
372         # Update proposed move (best so far, avoid timeout while find a
373             better move)
374         self.propose_move(best_move[0])
375         # print(move.i, move.j, value)

```

Code Listing 2: Improved Minimax agent from Assignment A3. `team33_A3/sudokuai.py`.

```

1  """
2  Competitive Sudoku AI.

4  Adapted from /naive_player/sudokuai.py

6  A1: iterative deepening minimax search with alpha beta pruning.
7  A2: heuristic search.
8  A3: heuristic search with unsolvable moves list.

10 """

12 from random import shuffle

14 # Import types and libraries
15 from competitive_sudoku.sudoku import GameState, Move, SudokuBoard
16 from competitive_sudoku.sudokuai import SudokuAI

18 from .utils import block_range, StateMatrixT, SavedDataT, is_unsolvable # Game
19                                     specific logic
20 from .utils import (
21     block_index,
22     calculate_filling_rate,
23     calculate_move_score,
24     is_illegal,
25     next_player,
26 )

28 class SudokuAI(SudokuAI):
29     """
30     Sudoku AI agent that computes a move for a given sudoku configuration.
31     """

33     def minimax(
34         self,
35         *,

```

```

36     game_state: GameState,
37     move: Move,
38     state_matrix: StateMatrixT,
39     current_player: int,
40     maximizing_player: int,
41     depth: int,
42     alpha: float = float("-inf"),
43     beta: float = float("inf"),
44 ) -> float:
45     """
46     Returns the score of a given move.
47
48     Minimax search that considers the perspectives of two players .
49     Searches until some depth before returning the score .
50     Uses alpha beta pruning to avoid searching branches which cannot lead to
51         better results .
52
53     :param game_state: The current game state. Describes the board, scores ,
54         unsolvable moves and move history.
55     :param move: The move to be evaluated
56     :param state_matrix: State matrix containing the initial set of moves and
57         whether they are still legal
58     :param current_player: The player who's turn it is . Will be the same
59         throughout the turn. (0 or 1 for first or
60         second player)
61     :param maximizing_player: The player who's score is to be maximised. (0 or 1
62         for first or second player)
63     :param depth: The maximum depth to search before returning the score .
64     :param alpha: The highest so far value for alpha beta pruning. ( initially  -inf
65         )
66     :param beta: The lowest so far value for alpha beta pruning. ( initially  inf)
67     :return: The score of the move (higher is better for maximizing player, lower
68         is better for minimizing player)
69
70     """
71
72     # Apply move, update resulting scores
73     # Update legal moves and count newly invalidated moves
74     game_state.board.put(move.i, move.j, move.value)
75     _score_achieved = calculate_move_score(game_state, move)
76     game_state.scores[current_player] += _score_achieved
77     _moves_invalidated = self.update_legal(game_state, move, state_matrix)
78
79     # Switches perspective to other player
80     current_player = next_player(current_player)
81
82     # Search until game is finished or maximum depth is reached
83     if depth == 0 or state_matrix["legal_count"] == 0:
84         # evaluate the current board
85         best_value = self.evaluate_state(
86             maximizing_player,
87             game_state,

```

```

79         )
80     else:
81         if maximizing_player == current_player: # maximising player
82             best_value = float("-inf")
83             for try_move in state_matrix["initial"]:
84                 if not state_matrix["legal"][try_move.i][try_move.j][try_move.
                    value]:
85                     continue
86                 # Recurse and find value up to some depth
87                 value = self.minimax(
88                     game_state=game_state,
89                     move=try_move,
90                     state_matrix=state_matrix,
91                     current_player=current_player,
92                     maximizing_player=maximizing_player,
93                     depth=depth - 1,
94                     alpha=alpha,
95                     beta=beta,
96                 )
97                 best_value = max(best_value, value)
98                 alpha = max(alpha, best_value)
99                 if beta < alpha:
100                     break
101         else: # minimising player
102             best_value = float("inf")
103             for try_move in state_matrix["initial"]:
104                 if not state_matrix["legal"][try_move.i][try_move.j][try_move.
                    value]:
105                     continue
106                 # Recurse and find value up to some depth
107                 value = self.minimax(
108                     game_state=game_state,
109                     move=try_move,
110                     state_matrix=state_matrix,
111                     current_player=current_player,
112                     maximizing_player=maximizing_player,
113                     depth=depth - 1,
114                     alpha=alpha,
115                     beta=beta,
116                 )
117                 best_value = min(best_value, value)
118                 beta = min(beta, best_value)
119                 if beta < alpha:
120                     break
121         # Undo move and its effects
122         current_player = next_player(current_player)
123         state_matrix["legal_count"] += len(_moves_invalidated)
124         for inv_move in _moves_invalidated:
125             state_matrix["legal"][inv_move.i][inv_move.j][inv_move.value] = True
126         game_state.scores[current_player] -= _score_achieved
127         game_state.board.put(move.i, move.j, 0)

```

```

129     # Recursion result
130     return best_value

132 def evaluate_state( self , maximizing_player: int , game_state: GameState) -> float:
133     """
134     Calculates the score of a given game state.

136     :param maximizing_player: The player whose score is to be maximised. (0 or 1
137     for first or second player)
137     :param game_state: The current game state.
138     :return: The score of the game state. (higher is better for maximizing player,
139     and vice versa)
139     """
140     return game_state.scores[maximizing_player] - game_state.scores[next_player(
141         maximizing_player)]

142 def update_legal( self , game_state: GameState, move: Move, moves: StateMatrixT) -
143     > list:
144     """
145     Update which moves are legal in the recursive minimax search.

146     :param game_state: The current game state. Describes the board, scores, taboo
147     moves and move history.
147     :param move: The move to be evaluated
148     :param moves: State matrix containing the initial set of moves and whether
149     they are still legal
149     :return: A list of moves that were invalidated by the move.
150     """

152     _moves_invalidated = []
153     for row in range(game_state.board.board_height()):
154         # move invalidates another move if not already illegal , avoid double
155         counting
156         if moves["legal"][row][move.j][move.value]:
157             _moves_invalidated.append(Move(row, move.j, move.value))
158             moves["legal"][row][move.j][move.value] = False # Set legal status

159     for column in range(game_state.board.board_width()):
160         if moves["legal"][move.i][column][move.value]:
161             _moves_invalidated.append(Move(move.i, column, move.value))
162             moves["legal"][move.i][column][move.value] = False

164     for row, col in block_range(row=move.i, col=move.j, board=game_state.board):
165         if moves["legal"][row][col][move.value]:
166             _moves_invalidated.append(Move(row, col, move.value))
167             moves["legal"][row][col][move.value] = False
168     moves["legal_count"] -= len(_moves_invalidated)
169     return _moves_invalidated

171 def find_initial_moves( self , game_state: GameState) -> list[Move]:

```

```

172     """
173     Find all possible moves for a given state.

175     :param game_state: The current game state.
176     :return: The list of legal initial moves.
177     """
178     board_size = game_state.board.board_width()

180     # Generate possible moves
181     initial_moves = [
182         Move(i, j, value)
183         for i in range(board_size)
184         for j in range(board_size)
185         for value in range(1, board_size + 1)
186         if not is_illegal (move=Move(i, j, value), state=game_state)
187     ]

189     # Shuffle moves to be less predictable
190     shuffle (initial_moves)

192     return initial_moves

194 def calculate_free_cells ( self , game_state: GameState) -> dict:
195     """
196     Calculate how many free cells there are in each region of the sudoku (row,
197         column, block), given a game state.

198     It is used to assign each move a priority level.

200     :return: A dictionary containing the number of free cells , for each region .
201             row: list containing a value for each row of the sudoku, corresponding to
202                  the number of free cells in the row
203             col: list containing a value for each column of the sudoku, corresponding
204                  to the number of free cells in the column
205             block: bidimensional list containing a value for each block of the sudoku,
206                   corresponding to the number of free cells
207                   in the block

208     """
209     board_size = game_state.board.board_height()

211     free_cells = {
212         "row": [],
213         "col": [],
214         "block": [],
215     }
216     # Calculate free cells in rows
217     for y in range(board_size):
218         empty_count = sum(
219             1
220             for x in range(board_size)
221             if game_state.board.get(y, x) == SudokuBoard.empty

```

```

218         )

220         free_cells ["row"].append(empty_count)

222     # Calculate free cells in columns
223     for x in range(board_size):
224         empty_count = sum(
225             1
226             for y in range(board_size)
227             if game_state.board.get(y, x) == SudokuBoard.empty
228         )

230         free_cells ["col"].append(empty_count)

232     # Calculate free cells in blocks
233     for block_i in range(board_size // game_state.board.region_height()):
234         empty_list = []
235         for block_j in range(board_size // game_state.board.region_width()):
236             empty_count = 0
237             for row, col in block_range(
238                 row=block_i * game_state.board.region_height(),
239                 col=block_j * game_state.board.region_width(),
240                 board=game_state.board,
241             ):
242                 if game_state.board.get(row, col) == SudokuBoard.empty:
243                     empty_count += 1
244             empty_list.append(empty_count)
245         free_cells ["block"].append(empty_list)

247     return free_cells

249     def order_moves(
250         self, board: SudokuBoard, moves_list: list [Move], free_cells : dict
251     ) -> tuple[ list [Move], list [Move]]:
252         """
253         Sort a list of moves by their priority . Priority is based on the number of
254         regions the move allows to complete.
255         Moves that complete three regions have top priority and are therefore placed
256         at the top of the list ,
257         followed by those that complete two regions and those that complete one.
258         Moves that do not complete any region have no priority , and are returned in a
259         separate list .

260         :param board: The current board.
261         :param moves_list: List of moves that have to be sorted.
262         :param free_cells: dictionary containing the number of free cells for each
263                             region. It's used to assign priority to
264                             moves.

265         :return: Two lists , one with the moves with priority , sorted, the other with
266                 the moves without priority .

```

```

263      """
264
265      # Assign priority to initial moves based on how many regions each move
completes
266      priority = {
267          0: [], # None priority
268          1: [],
269          2: [],
270          3: [],
271      }
272
273      for move in moves_list:
274          row = move.i
275          col = move.j
276          block_i, block_j = block_index(row, col, board)
277
278          # Count how many regions the move completes
279          # sum counts the number of True values in the list
280          regions_to_be_completed = sum([
281              free_cells["row"][row] == 1,
282              free_cells["col"][col] == 1,
283              free_cells["block"][block_i][block_j] == 1
284          ])
285
286          priority[regions_to_be_completed].append(move)
287
288      return priority[3] + priority[2] + priority[1], priority[0]
289
290      def initialize_stored_data ( self , lower_limit: float , upper_limit: float ) ->
          SavedDataT:
291          """
292          Initialize stored data.
293          :param lower_limit: lower limit of the range of filling rate
294          values in which the list of moves that make Sudoku
unsolvable is calculated .
295          :param upper_limit: upper limit of the range of filling rate
296          :return: initial values for the dictionary containing the stored data.
297          """
298          return {
299              "unsolvable_moves": [],
300              "check_unsolvable_range": (
301                  lower_limit ,
302                  upper_limit,
303              ), # Check if a move makes the sudoku unsolvable only after
# the board is 20% complete and before is 85% complete
304              "unsolvable_list_building_started ": False ,
305              "unsolvable_list_building_finished ": False ,
306          }
307
308
309      def calculate_unsolvable_moves(
310          self , game_state: GameState, moves: list[Move], data: SavedDataT

```

```

311 ) -> None:
312     """
313     Calculate the list of unsolvable moves, simulating each move in the list
314         moves and checking
315         if the new gamestate represents a Sudoku impossible to solve.

316     :moves: list of possible unsolvable moves.
317     :data: SavedDataT dictionary containing the list of unsolvable moves.
318     """
319     for move in moves:
320         game_state.board.put(move.i, move.j, move.value)
321         if is_unsolvable(game_state.board):
322             data["unsolvable_moves"].append(move)
323         game_state.board.put(move.i, move.j, SudokuBoard.empty)

325 def update_unsolvable_moves(self, data: dict, moves: StateMatrixT) -> None:
326     """
327     Update the list of unsolvable moves,
328     checking if they are still present in the list of legal moves
329     calculated at the beginning of each turn.

331     :data: dictionary containing the list of unsolvable moves
332     :moves: list of legal moves
333     """
334     updated_moves = []
335     for move in data["unsolvable_moves"]:
336         if move in moves["initial "]:
337             updated_moves.append(move)
338             moves["initial "].remove(move)

340     data["unsolvable_moves"] = updated_moves

342 def compute_best_move(self, game_state: GameState) -> None:
343     """
344     Computes the best move for the agent and proposes it.

346     It initially proposes a random move, then sorts the moves based on their
347         priority and proposes the one with the
348         highest priority.

349     It then searches for the best move with minimax search.

350     Since the turn time is not known it proposes the best move found so far by
351         iteratively deepening the search, starting
352         with depth 2.

353     There is no need to search with depth 1, as moves are already sorted by
354         priority.

355     If the best move found so far leads to a negative result (calculated by the
356         evaluation function),
357     it proposes a move that makes the sudoku impossible to solve, in order to pass
358         the turn without entering new values.

```



```

355     The list of moves that make Sudoku impossible is calculated when the board is
356         20% full,
357     and is saved by the game engine so that it is also available in the next
358         rounds.
359     At each turn the list must be updated, to ensure that all the moves on it are
360         still legal .
361     If the time limit of the turn does not allow the list of moves that invalidate
362         the sudoku to be calculated ,
363     the agent tries to calculate it again when the board is 30% full.
364     If even then the time available is not enough, it waits for the board to be 40
365         % full, and so on.
366     When the list is empty (all moves have been used or are no longer legal), it
367         is recalculated

368
369     :param game_state: The current game state. Describes the board, scores and
370         move history.
371
372     """
373     board_size = game_state.board.board_height()
374
375     # Generate possible moves
376     initial_moves = self.find_initial_moves(game_state)
377
378     # Propose a random move (initial, avoid timeout)
379     self.propose_move(initial_moves[0])
380
381     # Order initial moves by priority
382     free_cells = self.calculate_free_cells(game_state)
383     priority_moves, non_priority_moves = self.order_moves(game_state.board,
384         initial_moves, free_cells)
385
386     # Initial moves are sorted by priority | set(initial_moves) == set(
387         priority_moves + non_priority_moves)
388     initial_moves = priority_moves + non_priority_moves
389
390     # Propose move with the highest priority ( initial , avoid timeout)
391     self.propose_move(initial_moves[0])
392
393     # Initialize legal moves matrix with all moves being illegal
394     legal_moves = [
395         [[False for v in range(board_size + 1)] for col in range(board_size)]
396         for row in range(board_size)
397     ]
398
399     # Mark initial moves as legal in the matrix
400     for move in initial_moves:
401         legal_moves[move.i][move.j][move.value] = True
402
403     # Initialize state matrix
404     state_matrix: StateMatrixT = {
405         "initial": initial_moves,

```

```

396         "legal_count": len( initial_moves ),
397         "legal": legal_moves,
398     }

400     data: SavedDataT = self.load()
401     if data is None: # if no data is saved, initialize it with default values
402         data = self . initialize_stored_data (0.2, 0.85)
403         self . save(data)

405     # If the previous attempt to build unsolvable moves list failed ,
406     # then increase the lower limit of the filling range by 10% (this number was
         chosen arbitrarily )

407     if (
408         data[" unsolvable_list_building_started "] is True
409         and data[" unsolvable_list_building_finished "] is False
410     ):
411         data["check_unsolvable_range"] = (
412             data["check_unsolvable_range"][0] + 0.1,
413             data["check_unsolvable_range"][1],
414         )
415         data[" unsolvable_list_building_started "] = False

417         self . save(data)

419     filling_rate = calculate_filling_rate ( free_cells )
420     # First turn or unsolvable move list is empty (all unsolvable moves have been
         used or are not legal anymore)

421     if (
422         len(data["unsolvable_moves"]) == 0
423         and filling_rate > data["check_unsolvable_range"][0]
424         and filling_rate < data["check_unsolvable_range"][1]
425     ):
426         data[" unsolvable_list_building_started "] = True
427         self . save(data)

429     # Only moves that not complete any regions can invalidate the game
430     self . calculate_unsolvable_moves(game_state, non_priority_moves, data)

432     data[" unsolvable_list_building_finished "] = True
433     self . save(data)

435     # Check which unsolvable moves can still be used and remove them from moves
436     if len(data["unsolvable_moves"]) > 0:
437         self . update_unsolvable_moves(data, state_matrix)
438         self . save(data)

440     # Iteratively increase the search depth of minimax
441     best_move: tuple[Move, float] | None = None
442     for depth_limit in range(2, len(state_matrix[" initial "]), 1):
443         # Evaluate different moves based on minimax
444         for move in state_matrix[" initial "]:

```

```

445         # Player_index is 0 or 1 ( first or second player )
446         player_index = game_state.current_player() - 1

448         value = self.minimax(
449             game_state=game_state,
450             move=move,
451             state_matrix=state_matrix,
452             current_player=player_index, # Current player is self
453             maximizing_player=player_index, # Maximising own score
454             depth=depth_limit,
455         )

457         if best_move is None or value > best_move[1]:
458             best_move = (move, value)
459             # If every explored move leads to a negative score, propose a
460             unsolvable move
461             if best_move[1] < 0 and len(data["unsolvable_moves"]) > 0:
462                 self.propose_move(data["unsolvable_moves"][0])
463             else:
464                 self.propose_move(best_move[0])

```

Code Listing 3: MCTS agent, the second agent from Assignment A3.
team33_A3_MCTS/sudokuai.py.

```

1  """Competitive Sudoku AI.

3  Adapted from /naive_player/sudokuai.py

5  A1: iterative deepening minimax search with alpha beta pruning.
6  A2: heuristic search.
7  A3: different improvements and a second strategy.
8  """

10 from typing import Literal, TypedDict, TypeVar, Type, Union

12 from competitive_sudoku.sudoku import GameState, Move, SudokuBoard
13 from competitive_sudoku.sudokuai import SudokuAI
14 from random import shuffle, choice
15 from copy import deepcopy
16 from math import log

18 DEBUG = False

20 if DEBUG:
21     from graphviz import Digraph
22     import uuid

25 # Game specific logic not related to strategy stored in utils.py
26 from .utils import (
27     block_range,
28     calculate_move_score,

```

```

29     is_illegal ,
30     next_player,
31     PlayerID,
32 )

34 # Hyperparameters
35 MCTS_EXPLORATION = 2

38 def build_graph(
39     node: "MCGTNode", max_player: PlayerID, graph=None, node_name=None,
40     depth_limit=None
41 ):
42     """ Visualise the MCTS game tree using graphviz """
43     if node_name is None:
44         node_name = uuid.uuid4().hex

45     if graph is None:
46         graph = Digraph()
47         graph.node(name=node_name, label=str(node).replace(";", "\n"))

49     for child in node.children :
50         child_name = uuid.uuid4().hex

52         graph.node(
53             name=child_name,
54             label=str(child).replace(";", "\n"),
55             color="green" if child.player == max_player else "black",
56         )
57         graph.edge(tail_name=node_name, head_name=child_name)

59         if depth_limit is None or child.depth < depth_limit:
60             build_graph(child, max_player, graph, child_name, depth_limit)

62     return graph

65 class StateMatrixT(TypedDict):
66     """Type of state matrix used in MCTS"""

67     initial : list [Move]
68     legal : list [ list [ list [bool]]]
69     legal_moves_count: int

73 class MCGTNode:
74     """Monte Carlo Game Tree Node"""

75     def __init__(
76         self,
77         parent: Union["MCGTNode", None],

```

```

79         move: Move,
80         depth: int,
81         player: PlayerID,
82         visited=0,
83         score=0,
84     ):
85         """MC Game tree node
86         :param parent: parent node, allows backtracking
87         :param move: move that was made to reach this node
88         :param depth: depth of this node in the tree
89         :param visited: number of times this node has been visited
90         :param score: score of the game at this node
91         """
92         self.move = move
93         self.visited = visited
94         self.score = score
95         self.parent = parent
96         self.children: list[MCGTNode] = [] # Children are added in expansion phase
97         self.depth = depth
98         self.player: PlayerID = player

100     @property
101     def is_leaf(self):
102         return not self.children

104     @property
105     def average_score(self):
106         return self.score / self.visited if self.visited > 0 else 0

108     @property
109     def ucb(self):
110         """Upper Confidence Bound of this node"""
111         if self.visited == 0:
112             return float("inf")

114         return self.average_score + MCTS_EXPLORATION * (
115             (log(self.parent.visited) / self.visited) ** 0.5
116         )

118     def __str__(self):
119         return f"{str(self.move)}; n = {self.visited}; q = {self.score}"

122 class MCGameTree:
123     def __init__(
124         self,
125         root: MCGTNode,
126         game_state: GameState,
127         state_matrix: StateMatrixT,
128         max_player: PlayerID,
129     ):

```

```

130      """Monte Carlo Game Tree

132      :param root: root node of the tree
133      :param game_state: current game state
134      :param state_matrix: state matrix used to track legal moves
135      :param max_player: player to maximize score for
136      """
137      self.root = root
138      self.maximizing_player = max_player
139      self.game_state = game_state
140      self.state_matrix = state_matrix

142      def selection ( self , node: MCGTNode) -> MCGTNode:
143          """Select a leaf node to expand.

145          Recursively select the child with the highest UCB until a leaf node is reached
          .

147          :param node: node to start selection from.
148          :return: leaf node.
149          """
150          if node.is_leaf :
151              return node

153          # Remove illegal moves (can occur because of reloading the tree)
154          node.children = [
155              child for child in node.children
156              if self.state_matrix["legal"][child.move.i][child.move.j][child.move.value
157              ]

159          # Select child with highest UCB
160          return self.selection (max(node.children, key=lambda child: child.ucb))

162      def expansion( self , node: MCGTNode) -> MCGTNode:
163          """Expand a leaf node.

165          A node is expanded when it is selected twice.
166          Expansion adds all legal moves in the current game state as children of the
              node.

168          :param node: node to expand.
169          :return: node to simulate from.
170          """
171          # Apply all moves from node to the root (excluding)
172          self.apply_move_on_node(node)

174          # If v is a terminal state of the game, move to Backpropagation phase
175          if self.state_matrix["legal_moves_count"] == 0:
176              return node

```

```

178     # If n(v) = 0, move to Simulation phase with node v
179     if node.visited == 0:
180         return node

182     # If n(v) > 0, add new states reached from legal moves in v
183     # TODO use array of legal moves instead of matrix and initial moves?
184     for move in self.state_matrix["initial "]:
185         if self.state_matrix["legal "][move.i][move.j][move.value]:
186             node.children.append(
187                 MCGTNode(node, move, node.depth + 1, next_player(node.player))
188             )

190     return choice(node.children)

192 def simulation( self , node: MCGTNode) -> int:
193     """Simulate a game from a node.
194
195     Simulate a game from a node by applying random moves until the game ends.
196
197     :param node: node to simulate from.
198     :return: score of the game. 1 when the maximizing player wins, 0 otherwise.
199     """

201     self.apply_move_on_node(node)
202     player = next_player(node.player)
203     while self.state_matrix["legal_moves_count"] > 0: # until game ends
204         # Select a random move
205         # TODO use array of legal moves instead of matrix and initial moves?
206         move = None
207         while move is None:
208             move = choice(self.state_matrix["initial "])
209             if not self.state_matrix["legal "][move.i][move.j][move.value]:
210                 move = None
211         # Apply move
212         self.apply_move(player, move)
213         player = next_player(player)

215     return (
216         1
217         if self.game_state.scores[self.maximizing_player]
218         > self.game_state.scores[next_player(self.maximizing_player)]
219         else 0
220     )

222 def backpropagation( self , node: MCGTNode | None, score: int) -> None:
223     """Backpropagation the score from a node to the root
224
225     :param node: node to start backpropagation from
226     :param score: score to backpropagation
227     """
228     if node is None:

```

```

229         return

231     node.visited += 1
232     node.score += score if node.player == self.maximizing_player else -score
233     self.backpropagation(node.parent, score)

235     def find_best_child( self ) -> Move:
236         """Find the best child of the root node

238         Is based on the best average score of the children of the root node.

240         :return: best move found so far
241         """
242         best_child = None
243         best_avg_score = float("-inf")
244         for child in self.root.children :
245             if child.average_score > best_avg_score:
246                 best_avg_score = child.average_score
247                 best_child = child

249         if best_child is None: # Happens in the first iteration
250             return None

252         return best_child.move

254     def iterate ( self ) -> Move:
255         """Perform one iteration of MCTS.

257         Runs all phases of MCTS once.
258         Selection , expansion, simulation and backpropagation.

260         :return: best move found so far as defined in 'find_best_child'.
261         """
262         # TODO optimise
263         selected_leaf = self.selection ( self.root )
264         selected_leaf = self.expansion( selected_leaf )

266         score = self.simulation ( selected_leaf )
267         self.backpropagation( selected_leaf , score )

269         return self.find_best_child ()

271     def apply_move_on_node(self, node: MCGTNode) -> None:
272         """Recursively apply moves to the game state"""
273         # If node is root, no moves to apply (root is a dummy node)
274         if node is self.root:
275             return

277         # If node parent is not a root, apply parent's moves first
278         if node.parent is not None:
279             self.apply_move_on_node(node.parent)

```



```

281         self .apply_move(node.player, node.move)

283     def apply_move(self, player : PlayerID, move: Move):
284         self .game_state.board.put(move.i, move.j, move.value)
285         self ._apply_move_update_state_matrix(move)
286         self .game_state.scores[ player ] += calculate_move_score(self.game_state, move)

288     def _apply_move_update_state_matrix(self, move: Move):
289         """Update the state matrix when a move is made"""
290         # Update which moves are legal
291         changes = 0
292         for idx in range(self .game_state.board.N):
293             # move invalidates another move if not already illegal , avoids double
counting
294             if self .state_matrix[" legal "][ idx ][ move.j ][ move.value ]:
295                 self .state_matrix[" legal "][ idx ][ move.j ][ move.value ] = False
296                 changes += 1

298             if self .state_matrix[" legal "][ move.i ][ idx ][ move.value ]:
299                 self .state_matrix[" legal "][ move.i ][ idx ][ move.value ] = False
300                 changes += 1

302         for row, col in block_range(
303             row=move.i, col=move.j, board=self.game_state.board
304         ):
305             if self .state_matrix[" legal "][ row ][ col ][ move.value ]:
306                 self .state_matrix[" legal "][ row ][ col ][ move.value ] = False
307                 changes += 1

309         self .state_matrix[" legal_moves_count"] -= changes

312 class SudokuAI(SudokuAI):
313     """
314     Sudoku AI agent that computes a move for a given sudoku configuration .
315     """

317     def compute_best_move(self, game_state: GameState) -> None:
318         """Computes the best move for the agent and proposes it .
319
320         Will initially propose a random move, then evaluate different moves based on
minimax search.
321
322         Since the turn time is not known it will propose the best move found so far by
iteratively deepening the search .
323
324         :param game_state: The current game state. Describes the board, scores and
move history.
325
326         """
327         board_size = game_state.board.board_height()
328         num_blocks = board_size // game_state.board.region_height()

```

```

328     # TODO also save game state matrix?
329     # player_index is 0 or 1 ( first or second player)
330     player_index = game_state.current_player() - 1

332     # List of initially legal moves, used as subset to avoid iterating over all
           moves

333     initial_moves = [
334         Move(i, j, value)
335         for i in range(board_size)
336         for j in range(board_size)
337         for value in range(1, board_size + 1)
338         if not is_illegal (move=Move(i, j, value), state=game_state)
339     ]
340     # Propose a random move (initial, avoid timeout)
341     self.propose_move(choice(initial_moves))

343     # Avoid repeated regeneration of legal moves by tracking their status
344     legal_moves = [
345         [[False for _ in range(board_size + 1)] for _ in range(board_size)]
346         for _ in range(board_size)
347     ]
348     for move in initial_moves:
349         legal_moves[move.i][move.j][move.value] = True

351     state_matrix: StateMatrixT = {
352         "initial": initial_moves,
353         "legal": legal_moves,
354         "legal_moves_count": len(initial_moves),
355     }

357     cache = self.load()
358     if cache is None: # First turn
359         # Create dummy root node for MCTS, the move "belongs" to the opponent
360         dummy_root = MCGTNode(
361             None, Move(-1, -1, -1), 0, next_player(player_index)
362         )
363         game_tree = MCGameTree(
364             dummy_root,
365             game_state,
366             state_matrix,
367             player_index, # Maximizing player
368         )
369         cache = {}
370         cache["tree"] = game_tree
371         # cache["state_matrix"] = state_matrix
372         self.save(cache)
373     else: # Return with cached state
374         game_tree = cache["tree"]
375         # Go to new position in tree
376         moves = game_state.moves[-2:] # Moves made since last turn (from argument

```

```

377         node = game_tree.root
378     for move in moves:
379         if move not in game_state.taboo_moves: # This move was indeed made
380             since it is not taboo
381             # game_tree.apply_move(player_index, move) # Update state and
382             state matrix
383             for child in node.children: # Move to child
384                 if child.move == move:
385                     node = child
386                     node.parent = None
387                     break
388             else: # Child not found
389                 node = MCGTNode(None, move, node.depth + 1, next_player(
390                     node.player))
391
392     game_tree.root = node
393     self.save(cache)
394
395     # Remove data not used in MCTS
396     del game_state.taboo_moves
397     del game_state.moves
398     del game_state.initial_board
399
400     # Keep improving the estimate of best move using MCTS until timeout
401     while True:
402         game_tree.state_matrix = deepcopy(state_matrix)
403         game_tree.game_state = deepcopy(game_state)
404         best_move = game_tree.iterate()
405
406         if best_move is not None:
407             self.propose_move(best_move)
408             # Store game tree to reuse
409             # TODO save only if timeout is close
410             if game_tree.root.visited % 10 == 0:
411                 self.save(cache)
412
413         if DEBUG and game_tree.root.visited == 25:
414             graph = build_graph(game_tree.root, player_index, depth_limit=5)
415             graph.view()
416             print("Graph built")
417             exit()

```

Code Listing 4: Different utility functions for game-specific logic. team33_A3/utils.py and team33_MCTS/utils.py.

```

1  """
2  Module containing helper functions for the sudoku game.
3
4  These functions are used for game specific logic, such as checking if a move is legal
5  or calculating the score of a move.
6  They do not contain any strategic logic specific to the AI agent.
7  """

```

```

8  from typing import Iterator, TypedDict

10 from competitive_sudoku.sudoku import GameState, Move, SudokuBoard

13 class StateMatrixT(TypedDict):
14     """
15     Type definition for the state matrix.

17     :param initial: List of moves that were initially on the board
18     :param legal: Array of shape (board_size, board_size, board_size + 1)
19                   where legal[i, j, k] is True if Move(i, j, k) is legal
20     :param legal_count: count of legal moves
21     """
22     initial: list[Move]
23     legal: list[list[list[bool]]]
24     legal_count: int

27 class SavedDataT(TypedDict):
28     """
29     Type definition for the saved data.

31     :param unsolvable_moves: List of moves that make the board unsolvable
32     :param check_unsolvable_range: When to calculate the unsolvable moves
33     :param unsolvable_list_building_started: True if the unsolvable list building has
34                                             started
35     :param unsolvable_list_building_finished: True if the unsolvable list building has
36                                             finished
37     """
38     unsolvable_moves: list[Move]
39     check_unsolvable_range: tuple[float, float]
40     unsolvable_list_building_started: bool
41     unsolvable_list_building_finished: bool

43 def next_player(current_player: int) -> int:
44     """
45     Returns the next player.

47     :param current_player: the current player.
48     :return: the next player.
49     """
50     return (current_player + 1) % 2

53 def block_index(row: int, col: int, board) -> tuple[int, int]:
54     """
55     Transform in which block a certain coordinate is.

```

```

57     Enumerates blocks in the same way as coordinates, top left is (0, 0).

59     :param row: a row index inside the block.
60     :param col: a column index inside the block.
61     :param board: the board to which the indices belong.
62     :return: the indices of the block in the board ( vertical , horizontal ).
63     """
64     return (
65         row // board.region_height(), # floor division
66         col // board.region_width(),
67     )

70 def block_range(*, row: int, col: int, board: SudokuBoard) -> Iterator[tuple[int, int]]
71     """
72     Return the range of indices in a block.

74     :param row: a row index inside the block.
75     :param col: a column index inside the block.
76     :param board: the board to which the indices belong.
77     :return: an iterator of indices in the block.
78     """
79     region_width, region_height = board.region_width(), board.region_height()
80     block_indices = block_index(row, col, board)

82     region_start = {
83         "x": block_indices[1] * region_width,
84         "y": block_indices[0] * region_height,
85     }

87     region_end = {
88         "x": (block_indices[1] + 1) * region_width,
89         "y": (block_indices[0] + 1) * region_height,
90     }

92     for row in range(region_start["y"], region_end["y"]):
93         for col in range(region_start["x"], region_end["x"]):
94             yield row, col

97 def is_illegal (*, move: Move, state: GameState) -> bool:
98     """
99     Returns whether a move is illegal .

101     A move is illegal if it puts a duplicate value in a position , block, row or
102     column.
103     Additionally , moves should not be 'taboo', meaning that they make the board
104     unsolvable.
105     Illegal moves are not allowed and will result in a loss .

```

```

105     :param move: The move to be checked.
106     :param state: The current state of the game.
107     :return: Whether the move is illegal (True) or not (False).
108     """

110     # Check if square is empty
111     if state.board.get(move.i, move.j) != SudokuBoard.empty:
112         return True

114     # Check if move is taboo
115     if move in state.taboo_moves:
116         return True

118     # Check duplicate value in row
119     if any(
120         state.board.get(row, move.j) == move.value
121         for row in range(state.board.board_height())
122     ):
123         return True

125     # Check duplicate value in column
126     if any(
127         state.board.get(move.i, col) == move.value
128         for col in range(state.board.board_width())
129     ):
130         return True

132     # Lastly check for duplicate values in the region
133     return any(
134         state.board.get(row, col) == move.value
135         for row, col in block_range(row=move.i, col=move.j, board=state.board)
136     )

139 def calculate_move_score(game_state: GameState, move: Move) -> int:
140     """
141     Check if a move completes any regions and returns the score earned.

143     Static method, uses less memory since it does not need to be instantiated.

145     :param game_state: The current game state. Describes the board, scores and move
146         history.
147     :param move: The move to be evaluated
148     :return: The score earned by the move (0, 1, 3 or 7)
149     """
150     row_complete = col_complete = block_complete = True

151     # Check if completed a row
152     for col in range(game_state.board.board_width()):
153         if game_state.board.get(move.i, col) == SudokuBoard.empty:

```

```

154         row_complete = False
155         break

157     # Check if completed a column
158     for row in range(game_state.board.board_height()):
159         if game_state.board.get(row, move.j) == SudokuBoard.empty:
160             col_complete = False
161             break

163     # Check if completed a block
164     for row, col in block_range(row=move.i, col=move.j, board=game_state.board):
165         if game_state.board.get(row, col) == SudokuBoard.empty:
166             block_complete = False
167             break

169     # Return score by move
170     regions_complete = int(row_complete) + int(col_complete) + int(block_complete)
171     return {
172         0: 0,
173         1: 1,
174         2: 3,
175         3: 7,
176     }[regions_complete]

179 def calculate_filling_rate ( free_cells : dict ) -> float:
180     """
181     Calculate the percentage of occupied cells on the sudoku board.

183     :param free_cells: A dictionary containing the number of free cells , for each
184         region .
185     row: list containing a value for each row of the sudoku, corresponding to the
186         number of free cells in the row
187     col: list containing a value for each column of the sudoku, corresponding to
188         the number of free cells in the column
189     block: bidimensional list containing a value for each block of the sudoku,
190         corresponding to the number of free cells
191         in the block

192     """
193     num_cells = len( free_cells ["row"]) * len( free_cells ["col"])
194     num_empty_cells = 0
195     for empty_in_row in free_cells ["row"]:
196         num_empty_cells += empty_in_row

197     return (num_cells - num_empty_cells) / num_cells

199 def is_unsolvable (board: SudokuBoard) -> bool:
200     """
201     Determines whether a sudoku is impossible to solve , given a game state.
202     It is used to create the unsolvable move list .

```

```

201         :return: True if the sudoku does not have any solutions , False otherwise.

202     """
203     board_size = board.board_height()
204     num_blocks = board_size // board.region_width()

205     # Consider rows of each block
206     for row_index in range(board_size):
207         for block_ind in range(num_blocks):
208             # List containing the values not present in the condidered row of the
209             # block
210             missing_values = [val for val in range(1, board_size + 1)]

211             block_row_index = row_index // board.region_height()
212             # List containing the indexes of the rows in the block that are not
213             # considered
214             missing_row_indexes = [
215                 index
216                 for index in range(
217                     block_row_index * board.region_height(),
218                     block_row_index * board.region_height() + board.region_height(),
219                 )
220                 if index != row_index
221             ]
222             # List containing the indexes of the columns without values in the
223             # considered row of the block
224             missing_col_indexes = [
225                 index
226                 for index in range(
227                     block_ind * board.region_width(),
228                     block_ind * board.region_width() + board.region_width(),
229                 )
230             ]
231             # Build the lists of missing values and missing indexes
232             for col_index in range(
233                 block_ind * board.region_width(),
234                 block_ind * board.region_width() + board.region_width(),
235             ):
236                 value = board.get(row_index, col_index)
237                 if value != SudokuBoard.empty:
238                     missing_values.remove(value)
239                     missing_col_indexes.remove(col_index)

240             # If the row of the block is empty, then it can't invalidate the sudoku
241             if len(missing_values) == board.region_height():
242                 continue

243             # Check the impossibility conditions of sudoku
244             for missing_value in missing_values:
245                 # If the value missing in the block row is present in the same block,

```



```

248     # then that value is not invalidating the sudoku
249     found_in_block = False
250     for row, col in block_range(
251         row=row_index,
252         col=block_ind * board.region_width(),
253         board=board,
254     ):
255         if board.get(row, col) == missing_value:
256             found_in_block = True
257             break
258     # Check the next missing value
259     if found_in_block:
260         continue

262     # Check invalidating values in rows (the missing value must be present
        in all missing rows)

263     found_in_rows_count = 0
264     for missing_row_index in missing_row_indexes:
265         for col in range(board_size):
266             if board.get(missing_row_index, col) == missing_value:
267                 found_in_rows_count += 1
268     # Check the next missing value
269     if found_in_rows_count != len(missing_row_indexes):
270         continue

272     # Check invalidating values in columns (the missing value must be
        present in all missing columns)

273     found_in_cols_count = 0
274     for missing_col_index in missing_col_indexes:
275         for row in range(board_size):
276             if board.get(row, missing_col_index) == missing_value:
277                 found_in_cols_count += 1
278     # Check the next missing value
279     if found_in_cols_count != len(missing_col_indexes):
280         continue

282     # If the value has been found in all rows and columns and is not in
        the block,
283     # then the sudoku is impossible to solve
284     return True

286     # All the missing values have been checked and none of them is invalidating the
        sudoku

287     return False

```
